Chapter - 1
Learning From Data

Aviral Janveja

1 Introduction

Let us look at the fundamentals of machine learning, utilizing an example from the domain of finance. Suppose someone applies for a bank loan. Now, the bank must decide whether granting the loan is a good idea or not, since their goal is to maximize profit while minimizing risk. Naturally, the bank has no magic formula to predict whether a customer will make timely repayments or not. However, it is indeed possible to look at similar customers from the past and examine how their repayment behavior turned out to be.

The **key idea** here, is to use past data points to uncover an underlying pattern that connects customer information with their repayment behavior. So that, we can use it to predict the creditworthiness of fresh applicants. Using this example, let us now outline the main elements of the machine learning process :

1.1 The Underlying Pattern

As discussed above, in machine learning, the idea is to uncover the underlying pattern from a given set of data points, such that we can use it to predict or generate future observations. In the bank loan example, this underlying pattern is the relation that connects customer information with their repayment behavior. Although we do not see this relation directly, we do get glimpses of it through the available data points.

In this section, we will describe this underlying pattern mathematically. But first, let us first consider a simpler scenario. Suppose a bank employs the following rule for loan approval :

> Approve every applicant whose annual income exceeds Rs. 7 lakh, otherwise reject the application.

This is a simple rule with no uncertainty. Every input produces exactly one correct output. Such a relation can be easily represented via a regular mathematical function. Let us call it the **target function** :

$$f(\mathbf{x})$$

For now, let us consider income as the only input metric for simplicity. So, if x = Rs. 9 lakhs per annum, then :

$$f(9) = +1$$

Indicating that the loan shall be approved, as per the above rule. However, we know that real world decision making is rarely this simple. Two applicants may have identical ages, salaries and employment histories, yet one eventually repays the loan while the other defaults. This can happen due to factors not captured in our data, such as job loss, unexpected medical expenses or simply the inherent uncertainty in human life. This leads us to an important observation :

> Most real world scenarios that we will deal with in machine learning are probabilistic in nature.

Hence, instead of fixed outputs, we get probabilities of different possible outputs, for a given input. For example :

$$P(y = +1 \mid x = 9) = 0.80$$
$$P(y = -1 \mid x = 9) = 0.20$$

Meaning that an applicant with a salary of Rs. 9 lakhs per annum, has an 80% chance of repaying the loan and a 20% chance of defaulting.

Therefore, instead of assuming a fixed output for every input, we describe the underlying pattern in terms of probabilities of different possible outputs, for a given input. Mathematically, this is represented by a conditional probability distribution, formally called the **target distribution** :

$$P(y \mid \mathbf{x})$$

It stands for "probability of y given $\mathbf{x}$" and is basically a function which gives us the probabilities of all possible outputs, for a given input, as shown in the above example.

Finally, it is important to **note** that the target function is simply a **special case** of the more general concept of target distribution. For example, when the probability of a certain output is 1 and thereby 0 for all others, the target distribution collapses to a target function and we get a single fixed output :

$$P(y \mid \mathbf{x}) = \begin{cases} 1 & \text{if } y = f(\mathbf{x}) \\ 0 & \text{if } y \neq f(\mathbf{x}) \end{cases}$$

1.2 The Data Set

The data set is simply a collection of data points generated by the underlying probabilistic process, as described above. In the loan approval example, each data point corresponds to a past customer and collecting N such data points gives us the required data set :

$$\mathbf{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), (\mathbf{x}_3, y_3)...(\mathbf{x}_N, y_N)\}$$

Where $\mathbf{x}_n$ is the application information of the nth customer, represented by an array of numbers as shown below :

$$\mathbf{x}_n = \begin{bmatrix} x_1 \\ x_2 \\ . \\ . \\ x_d \end{bmatrix}$$

Here x_1, x_2 and so on upto x_d are data attributes of a customer like age, annual income, years in residence, outstanding debt and so on. While the corresponding output :

$$y_n \in \{+1, -1\}$$

represents whether that customer eventually turned out to be creditworthy or not. Now, the data set does not reveal the target distribution completely, but it provides us with examples using which we can attempt to approximate it. In other words, the data set acts as our limited window into the underlying pattern.

2 Our First Learning Model : The Perceptron

So far, we have examined the two fundamental ingredients of every machine learning process, the **target distribution** and the **data set**. The next question is :

> How do we use these data points to learn something useful about the target distribution?

To answer this question, we introduce perceptron, a simple linear model. Let us suppose that for a particular customer :

$$P(y = +1 \mid \mathbf{x}_n) = 0.80$$
$$P(y = -1 \mid \mathbf{x}_n) = 0.20$$

The perceptron does not attempt to estimate these probabilities. Instead, it only predicts the final binary output. In this case, it would simply predict :

$$y = +1$$

Although, this is only a simplified representation of the underlying probabilistic pattern, it provides an excellent starting point for understanding how machine learning models work. More advanced learning models introduced later will attempt to capture the probabilistic nature of the target distribution more explicitly.

2.1 Hypothesis Function

The perceptron works by assigning a numerical weight to each input data attribute, such as age, income and so on. The weighted attributes are then summed together to produce a single score :

$$w_1 x_1 + w_2 x_2 + ... + w_d x_d$$

This weighted sum can be interpreted as a credit score assigned by the perceptron. The idea is pretty simple, if the credit score is greater than a certain threshold, the loan is approved, otherwise it is denied :

$$\begin{aligned} w_1x_1 + w_2x_2 + ... + w_dx_d > \text{threshold} \quad \text{(Approve)} \\ w_1x_1 + w_2x_2 + ... + w_dx_d \leq \text{threshold} \quad \text{(Deny)} \end{aligned}$$

For mathematical convenience, we rewrite the threshold as another weight :

$$w_0 = -\text{threshold}$$

along with an artificial attribute to go along with it :

$$x_0 = 1$$

Hence, the decision rule now becomes :

$$\begin{aligned} w_0x_0 + w_1x_1 + w_2x_2 + ... + w_dx_d > 0 \quad \text{(Approve)} \\ w_0x_0 + w_1x_1 + w_2x_2 + ... + w_dx_d \leq 0 \quad \text{(Deny)} \end{aligned}$$

Next, we write the above expression in vector notation to make it clean and concise :

$$w_0x_0 + w_1x_1 + .. + w_dx_d \; = \begin{bmatrix} w_0 & w_1 & . & . & w_d \end{bmatrix} \begin{bmatrix} x_0 \\ x_1 \\ . \\ . \\ x_d \end{bmatrix} \; = \; \mathbf{w^T x}$$

Since, the above formula outputs a real number that could be greater than, equal to or less than zero whereas our desired output is binary, we pass this value through the **sign** function, which outputs $+1$ if the value is greater than zero and -1 otherwise, hence giving us the final perceptron hypothesis function :

$$h(\mathbf{x}) = \text{sign}(\mathbf{w^T x})$$

Hypothesis Set

The mathematical form of the hypothesis function is now fixed. The only unknown quantities are the weights contained in the vector $\mathbf{w}$. Weights are real numbers and can therefore take infinitely many different values.

Each different choice of the weight vector thereby produces a different hypothesis. The collection of all such candidate hypotheses is called the hypothesis set, denoted by :

$$\mathbf{H} = \{h_1, h_2, h_3....\}$$

2.2 Perceptron Learning Algorithm

Having fixed the hypothesis function, the remaining task is to determine the values of the weight vector :

$$\mathbf{w}$$

We begin by initializing the weights randomly, usually to zero. The perceptron learning algorithm then iterates through the data set, searching for **errors**. An error on point $(\mathbf{x}_n, y_n)$ is when our hypothesis disagrees with the actual observed output :

$$\text{sign}(\mathbf{w}^\mathbf{T}\mathbf{x}_n) \neq y_n$$

Whenever such a point is found, the weight vector is updated by applying the following **update rule** :

$$\mathbf{w}_{new} = \mathbf{w} + y_n\mathbf{x}_n$$

The algorithm continues making such updates until no misclassified points remain in the training data or another stopping criterion is reached. And that is it, we then have our required set of final weights. The final weight vector determines the final hypothesis :

$$h_f(\mathbf{x}) = \text{sign}(\mathbf{w}^\mathbf{T}\mathbf{x})$$

which can then be used to classify future unseen examples.

2.3 Update Rule Justification

Let us examine how the above update rule actually nudges the perceptron in the right direction at each misclassified point. If a point $(\mathbf{x}_n, y_n)$ is misclassified, it means :

$$\text{sign}(\mathbf{w}^\mathbf{T}\mathbf{x}_n) \neq y_n$$

Accordingly, the perceptron updates the weight vector as follows :

$$\mathbf{w}_{new} = \mathbf{w} + y_n\mathbf{x}_n$$

Substituting the updated weight vector from above and ignoring the sign function for now, we get the following :

$$\mathbf{w}_{new}^\mathbf{T}\mathbf{x}_n = (\mathbf{w} + y_n\mathbf{x}_n)^\mathbf{T}\mathbf{x}_n$$

Expanding the right-hand side, we get :

$$\mathbf{w}_{new}^\mathbf{T}\mathbf{x}_n = \mathbf{w}^\mathbf{T}\mathbf{x}_n + y_n\mathbf{x}_n^\mathbf{T}\mathbf{x}_n$$

This is equivalent to :

$$\mathbf{w}_{new}^\mathbf{T}\mathbf{x}_n = \mathbf{w}^\mathbf{T}\mathbf{x}_n + y_n||\mathbf{x}_n||^2$$

Therefore, for a misclassified point, where $y_n = +1$ and $\text{sign}(\mathbf{w}^\mathbf{T}\mathbf{x}_n) = -1$, it implies $\mathbf{w}^\mathbf{T}\mathbf{x}_n \leq 0$ and after the update we get :

$$\mathbf{w}_{new}^{\mathbf{T}}\mathbf{x}_n = \mathbf{w}^{\mathbf{T}}\mathbf{x}_n + ||\mathbf{x}_n||^2$$

The update is adding some positive value to $\mathbf{w}^{\mathbf{T}}\mathbf{x}_n$ which might make $\mathbf{w}^{\mathbf{T}}\mathbf{x}_n > 0$ as required. Although, it does not guarantee an immediate sign correction, it clearly nudges the classifier in the right direction.

Similarly, when $y_n = -1$ and $\text{sign}(\mathbf{w}^{\mathbf{T}}\mathbf{x}_n) = +1$, implying $\mathbf{w}^{\mathbf{T}}\mathbf{x}_n > 0$, we get :

$$\mathbf{w}_{new}^{\mathbf{T}}\mathbf{x}_n = \mathbf{w}^{\mathbf{T}}\mathbf{x}_n - ||\mathbf{x}_n||^2$$

Hence, showing that each update to the weight vector nudges the perceptron in the required direction, improving its performance on that particular data-point. The following image illustrates the above discussion :

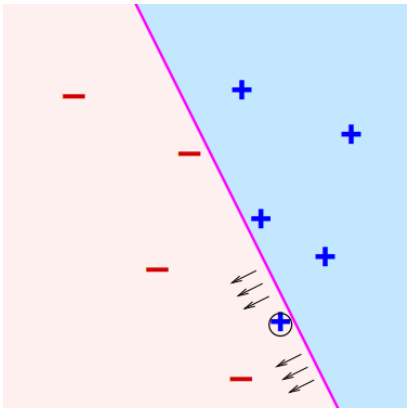


Figure 1: Perceptron

The pink line represents our linear classifier that is uniquely determined by our choice of weights. We have a misclassified **plus** point in the **minus** region and the update rule tries to correct this by nudging the classifier in the required direction.

Geometric Interpretation

Convergence Theorem

The convergence theorem for perceptron guarantees that if the data is **linearly separable**, repeated updates will eventually classify all points correctly.

Data being linearly separable simply means that, when the data points are plotted on a graph, it is possible to draw a straight-line that completely separates them into two distinct categories. The data in the above image for example, is linearly separable.

3 Types of Learning

As discussed in the introduction, machine learning is about learning patterns from data. Alternatively, we can put it as **using a set of observations to uncover an underlying process**. In either case, we have established that having a set of observations or data-points is a non-negotiable. However, the form in which the data is available to us can vary. Depending on this form and the nature of the learning task, machine learning is commonly divided into three main paradigms: supervised learning, unsupervised learning, and reinforcement learning.

3.1 Supervised Learning

In supervised learning, the model is trained on labeled data, where each example consists of : (**input, correct output**). The goal is to approximate a function from inputs to outputs that generalizes well to unseen data. For instance, the perceptron model above.

3.2 Unsupervised Learning

In unsupervised learning, the data is in the form of inputs only, without associated labels : (**input, no output**). The task is to discover hidden structures, patterns or clusters within the data. Examples include clustering, dimensionality reduction, and generative modeling.

3.3 Reinforcement Learning

In reinforcement learning, the model interacts with the environment by taking actions and receiving feedback in the form of rewards or penalties : (**action, feedback**). The model learns by trial and error to maximize its cumulative reward, gradually improving its decision-making. This setup is especially suited to sequential problems such as game playing, robotics and control systems.

4 References

1. CalTech Machine Learning Course - CS156, Lecture 1.
2. CalTech Machine Learning Course - CS156, Lecture 4.